



Figure2: An illustration of the architecture of our **CNN**, explicitly showing the delineation of responsibilities between the two **GPUs**. One GPU runs the layer-part at the top of the figure while the other runs the layer-parts at the bottom. The **GPUs** communicate only at certain layers. The network's input is 150.258-dimensional, and the number of neurons in the network's remaining layers is given by 253.440_186.624_64.896_43.264_4096_4096_1000.